

The background of the slide features a stylized representation of the Wordle logo, consisting of a 5x5 grid of colored squares. The colors include light green, light yellow, and light grey, arranged in a pattern that mimics the Wordle logo's design.

Wordle

Solver using Reinforcement Learning, Clustering and Search Algorithms

50.021 Artificial Intelligence

Group 3

Lee Min Shuen

Samuel Sim Wei Xuan

Riley Riemann Chin

Tan Hong Yew

Agenda

- Motivation
- Task Description
- Dataset
- Models
- Evaluations
- Graphical User Interfaces
- Improvements

Wordle

A 5x5 grid of squares in various shades of green, yellow, and grey, serving as a background for the word 'Motivation'.

Motivation

Motivation

- Addictive games (e.g. Flappy bird in 2014)




- Most recent craze is Wordle

Wordle

Motivation

- Famous math YouTube channel, 3Blue1Brown aimed to find the best starting word using information theory



Screenshot of Solving Wordle using information theory by 3Blue1Brown

- How might we apply AI techniques to solve Wordle?

Wordle

Task Description



How to play Wordle

Guess the **WORDLE** in six tries.

Each guess must be a valid five-letter word. Hit the enter button to submit.

After each guess, the color of the tiles will change to show how close your guess was to the word.

Examples

W E A R Y

The letter **W** is in the word and in the correct spot.

P **I** L L S

The letter **I** is in the word but in the wrong spot.

V A G **U** E

The letter **U** is not in the word in any spot.

SETTINGS



Hard Mode

Any revealed hints must be used in subsequent guesses




Task Description

- To use AI techniques to help solve the daily goal word in hard mode
- **Proposed algorithms:**
 - Reinforcement learning
 - Clustering algorithms
 - Search algorithms

General Strategy

Same starting word for all algorithms: CRANE

- Best starting word using information theory
- According to 3Blue1Brown



Screenshot of best opener for Wordle using information theory by 3Blue1Brown


Wordle

A 5x4 grid of colored squares representing Wordle feedback. The colors are: Row 1: Light Gray, Yellow, Light Gray, Yellow, Yellow; Row 2: Green, Yellow, Yellow, Yellow, Yellow; Row 3: Green, Yellow, Light Gray, Yellow, Green; Row 4: Green, Green, Green, Green, Green.

Dataset

Considerations

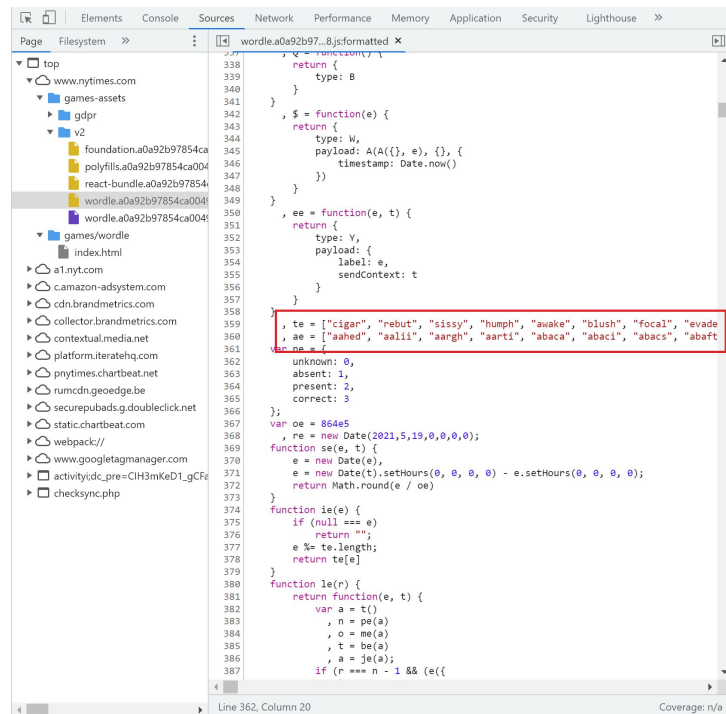
- 5-letter English words
- The Free Dictionary lists more than 150,000 5-letter English words



<https://www.thefreedictionary.com/5-letter-words.html>

Data Extraction

- Extracted 2 list of 5-letter words from Wordle javascript code
 - List 1: 2309 goal words
 - List 2: 12974 accepted words
- Training data: List 1 + List 2
 - 15283 words
- Testing data: List 1
 - 2309 goal words



```
...  
    return {  
      type: B  
    }  
  },  
  $ = function(e) {  
    return {  
      type: W,  
      payload: A(A(i), e), {}, {  
        timestamp: Date.now()  
      }  
    }  
  },  
  ee = function(e, t) {  
    return {  
      type: Y,  
      payload: {  
        label: e,  
        sendContext: t  
      }  
    }  
  },  
  te = ["cigar", "rebut", "sissy", "humph", "awake", "blush", "focal", "evade",  
  ae = ["aahed", "aalii", "aargh", "aarti", "abaca", "abaci", "abacs", "abeft",  
  ...  
  unknown: 0,  
  absent: 1,  
  present: 2,  
  correct: 3  
};  
var oe = 864e5  
var re = new Date(2021,5,19,0,0,0,0);  
function se(e, t) {  
  e = new Date(e),  
  e = new Date(t).setHours(0, 0, 0, 0) - e.setHours(0, 0, 0, 0);  
  return Math.round(e / oe)  
}  
function ie(e) {  
  if (null === e)  
    return "";  
  e = te.length;  
  return te[e]  
}  
function le(r) {  
  return function(e, t) {  
    var a = t(),  
    n = pe(a),  
    o = me(a),  
    t = be(a),  
    a = je(a);  
    if (r === n - 1 && (e({  
  ...  
}
```

Javascript code of Wordle from

<https://www.nytimes.com/games/wordle/index.html>

Some samples

- aahed
- aalii
- aargh
- aarti

- cigar
- rebut
- sissy
- humph

Wordle


A 5x5 grid of squares in shades of green, yellow, and grey, serving as a background for the word 'Models'.

Models

Quick Overview



**Reinforcement
Learning**




**Search
Algorithm**

Reinforcement Learning

3 Variations

- Q-Learning
- Q-Learning with Hierarchical Clustering
- Modified Q-Learning with Hierarchical Clustering



1) Q-learning: Q-table

Q-table: $Q(s,a)$ consists of

- Actions, a : Next word guess from the corpus
- States, s : List of words from the entire corpus of 15283 words

1) Q-learning: Q-function

Q-learning (Bellman equation):

$$Q(s, a) \leftarrow Q(s, a) + \alpha [R(s, a) + \gamma \max_{a'} Q'(s', a') - Q(s, a)]$$

The diagram illustrates the Q-learning Bellman equation with labels for each term:

- Reward from taking that action at the state**: Points to $R(s, a)$.
- Current Q-value**: Points to the first $Q(s, a)$ on the left.
- Learning rate**: Points to α .
- Discount factor**: Points to γ .
- Maximum expected future rewards across all possible actions and new state**: Points to $\max_{a'} Q'(s', a')$.
- Current Q-value**: Points to the $Q(s, a)$ at the end of the equation.

1) Q-learning: Reward System

Action:



Reward:


+10 +5 +5 +10 -1 = +24

- Higher score assigned to green letters than yellow letters
- Negative score assigned to black letters as it wastes a guess
- Reward for playing this action is updated and recorded in Q-table

1) Q-learning: Epsilon-Greedy

Choose an action at random instead of exploiting previous knowledge through the Q-table

- ϵ : probability of selecting a random action and exploring



1) Q-learning: Decaying Epsilon

Given that for each iteration, the goal is to guess the goal word within 6 tries

- Want to explore in first few tries
- Start to exploit only in later tries

Solution: Decaying epsilon using an inverse square function

$$\epsilon \leftarrow \frac{\epsilon}{(\text{stepnumber})^2}$$

1) Q-learning: Filtering

Hard-mode: Next guess has to use the color hints from the previous guess

Solution: Words no longer feasible are removed from the search space

- Green letters not at current position
- Yellow letters in current position
- Black letters present



1) Q-learning: Limitations

- Huge state space/Q-table: 15283×15283
- After each iteration of Wordle, Q-table is re-initialized to 0
- No carrying over of Q-table values, due to the filtering of state space

Solution: 2nd Variation - Addition of clustering algorithm


2) Q-learning with Hierarchical Clustering

Q-table: $Q(s,a)$ consists of

- Actions, a : Next cluster to select next word from
- States, s : List of cluster number

Clustering

- Reduces the size of state space
- Able to learn from previous iterations Q-tables
- Constant Q-table size, not affected by filtering



Why Hierarchical Clustering

- Sentiment meaning of words have 0 effect on guesses
- Tree structure more logical than a word vector space
- Bottom-up approach of grouping similar words together

Solution: Agglomerative hierarchical clustering + Levenshtein distance measure

Levenshtein distance

- Minimum number of single-character edits required to change one word to another



Example: Levenshtein distance = 3

2) Q-learning with Hierarchical Clustering

Everything else remains the same

- Reward system
- Epsilon-greedy policy
- Decaying epsilon
- Filtering function on corpus

3) Modified Q-learning with Hierarchical Clustering

- Since the daily wordle only comes from the 2309 goal words
- Solving this specific game only, and not a general 5-letter guesser

Modification:

- Training data: List 1 + ~~List 2~~
 - Only train on the 2309 goal words
- Testing data: List 1
 - 2309 goal words


Reinforcement Learning Strategy

1. Initialize a *Wordle* class which creates the Q-table and randomly selects a daily wordle from the list of goal words
 - a. Initialize a *Clustering* class to cluster on the corpus (for clustering version)
2. Uses the Q-learner to select the next guess
3. Passes the guess to an *Evaluation* class which gets the scoring and does the filtering
4. *Evaluation* class then passes the reward to update the Q-table for the next guess
5. Each run terminates when the goal word is reached
6. Reinitialize Q-table to 0 (for non-clustering version) or pass the Q-table to next run (for clustering version)

Search Algorithm

2 Variations

- Greedy-search
- Modified Greedy-search




1) Greedy Search

- Actions:
 - List of words from the entire corpus of 15283 words
- States:
 - Next guess from the same corpus
- Scoring:

Score each letter of the alphabet by frequency of occurrence in the corpus



Score each word by summing up the scores of each of its letters




1) Greedy Search

- Evaluation:
 - Evaluate green → yellow → grey letters to choose
 - Node with the lowest heuristic
 - Closest to the presumed goal state




Common limitations of Greedy Search

- Not optimal
 - Looks for the optimal local solution at every step
 - Cannot guarantee finding an optimal global solution
- Not complete
 - Might get stuck in a loop



1) Greedy Search: Limitations

- Not optimal
 - Cannot guarantee finding the goal word within the 6 tries
- ~~Not complete~~
 - Due to the nature of Wordle, the algorithm will always run to completion and will not get stuck in any loops.



2) Modified Greedy Search

Modification:

Similar to the reinforcement learning case

- Training data: List 1 + ~~List 2~~
 - Only train on the 2309 goal words
- Testing data: List 1
 - 2309 goal words



Search Algorithm Strategy


1. Randomly selects a daily wordle from the list of goal words
2. Initialize our word dictionary based on occurrence of words
3. Evaluates guess result and selects best option
4. Repeat step 3 until goal word is reached for termination

Summary of our 5 Models



Reinforcement Learning

- Q-learning
- Q-learning + hierarchical clustering
- Modified Q-learning + hierarchical clustering



Search Algorithm

- Greedy search
- Modified Greedy search

Wordle



Evaluations

Overview of Evaluations

1. Grid Search

- To find optimal hyper-parameters

2. Simulation analysis

- To analyze performance of models

Grid Search (Reinforcement Learning)


Different combinations of hyper-parameters

- Learning rate: [0.1, 0.01, 0.001]
 - Typical learning rates used
- Discounting factor: [0.5, 0.6, 0.7, 0.8, 0.9]
 - Our RL model is more focused on earlier rewards (from earlier guesses)
 - Instead of typical values of 0.99 or 0.9, lower discount factors are better
- Exploration rate: [0.5, 0.6, 0.7, 0.8, 0.9]
 - Due to decaying exploration rate, the lower bound cannot be too low
- Number of clusters: [6, 7, 8, 9, 10]
 - Arbitrary choice



Grid Search (Reinforcement Learning)

- For each of the RL variations:
 - Ran 100 simulations (runs of Wordle) for each hyperparameter combination
- Evaluate on 3 performance metrics
 - Win rate (Percentage of attempts with < 7 guesses)
 - Average number of guesses
 - Time taken for each Wordle run



Grid Search (Reinforcement Learning)

- Results of grid search, best combinations of hyperparameters

	Q-learning	Q-learning + Clustering	Modified Q-learning + Clustering
Learning rate	0.1	0.1	0.001
Discounting factor	0.8	0.9	0.9
Exploration rate	0.8	0.5	0.9
Number of clusters	-	6	9

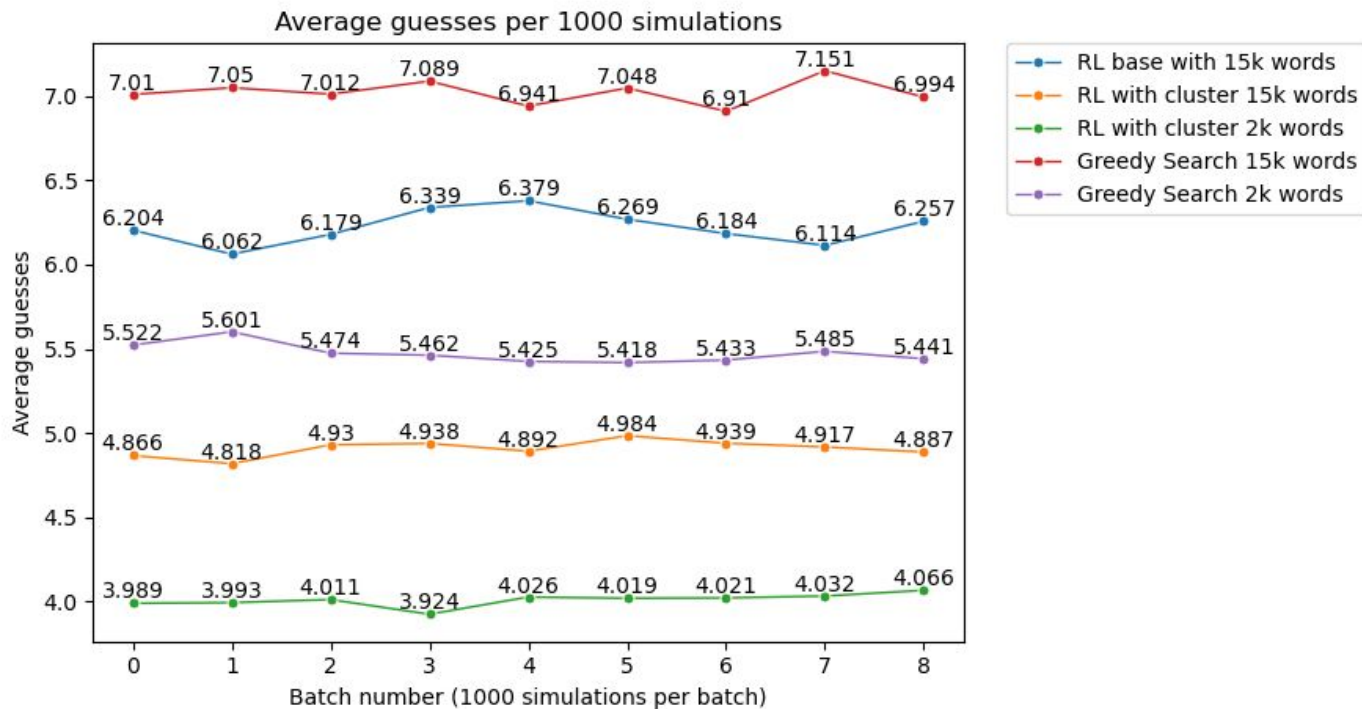
Simulation Analysis

With the found hyperparameters, to compare model performances

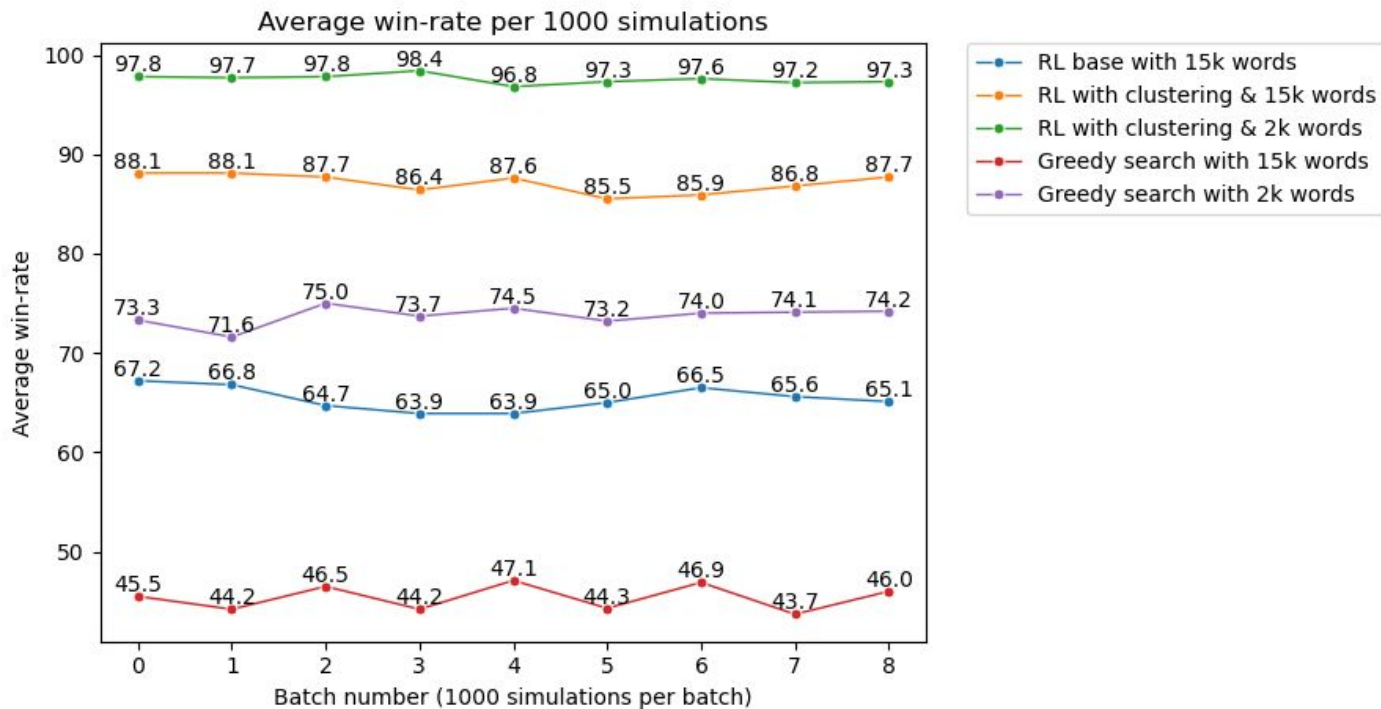
For each of the 5 models:

- Run 10,000 simulations
- Did a batch-means (size of 1000) of the number of guesses and win-rate
- Discard the first batch, to remove any initialisation bias
- Estimate mean and variance of each performance metric
- Compute the 95% confidence intervals

Results: Average number of guesses



Results: Average win-rate



Results: Confidence Intervals

	95% CI of average win-rate/%	95% CI of average number of guesses
Q-learning	(64.4, 66.5)	(6.213, 6.230)
Q-learning + Clustering	(86.4, 87.8)	(4.906, 4.910)
Modified Q-learning + Clustering	(97.4, 97.7)	(4.008, 4.010)
Greedy Search	(44.1, 46.6)	(7.019, 7.027)
Modified Greedy Search	(73.0, 74.4)	(5.471, 5.476)

Summary of Findings



Reinforcement Learning

- Requires less guesses
- Higher win-rate
- Clustering improves its performance



Search Algorithm

- Takes less time to run


Training on only the goal words gives consistently better performance for both

Best Model



Reinforcement Learning

- Q-learning
- Q-learning + hierarchical clustering
- **Modified Q-learning + hierarchical clustering**



Search Algorithm

- Greedy search
- Modified Greedy search

To be used for our daily wordle solver GUI

Wordle

Graphical User Interfaces



Overview of GUIs

1. Kivy GUI

- a. Visualize differences in performance of different models

2. PyGame GUI

- a. Observe an agent solving the daily Wordle

1) Kivy GUI

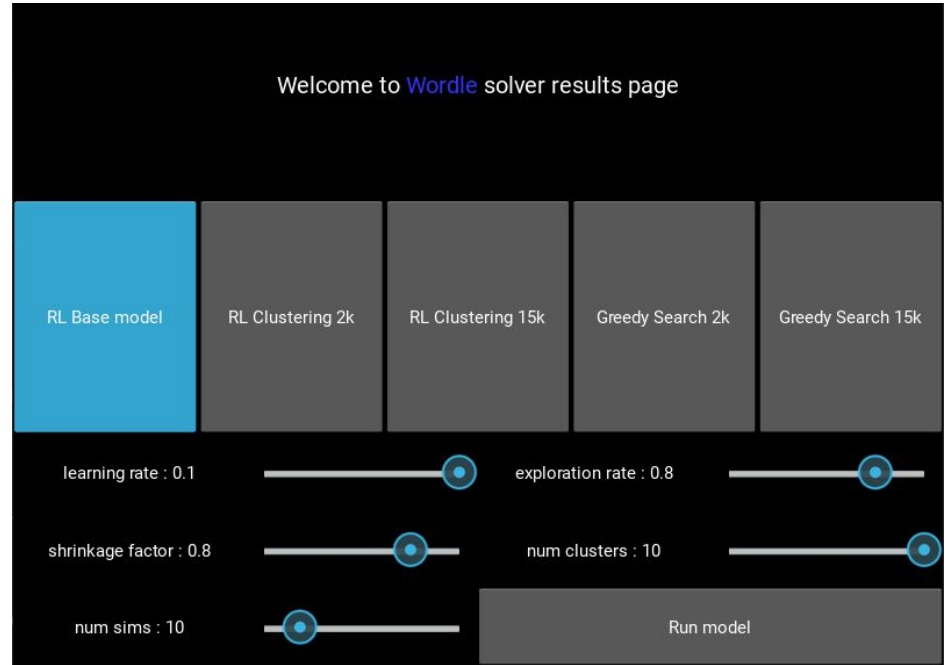
Visualize the difference in performances of different models

- Model selection
- Custom parameter values
- Graphical visualisation

1) Kivy GUI: Model Selection

Select from the five models:

- Q-Learning
- Q-Learning + Clustering
- Modified Q-Learning + Clustering
- Greedy search
- Modified Greedy search

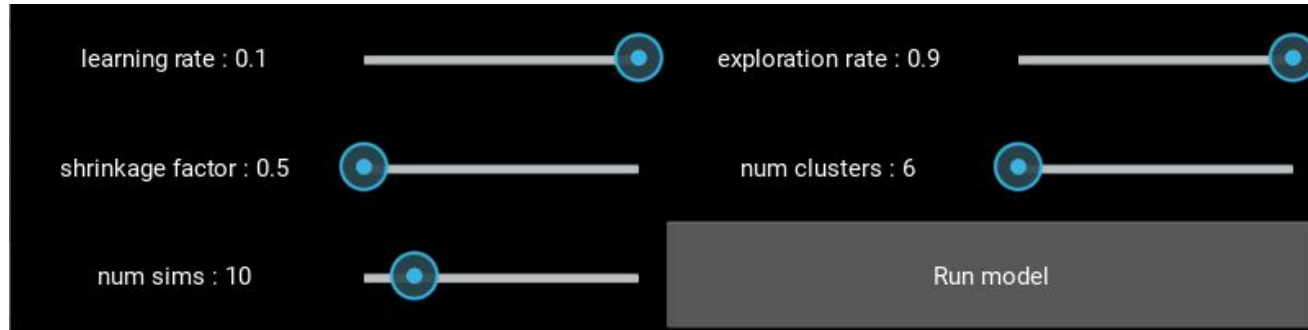


Screenshot of loading page of Kivy GUI

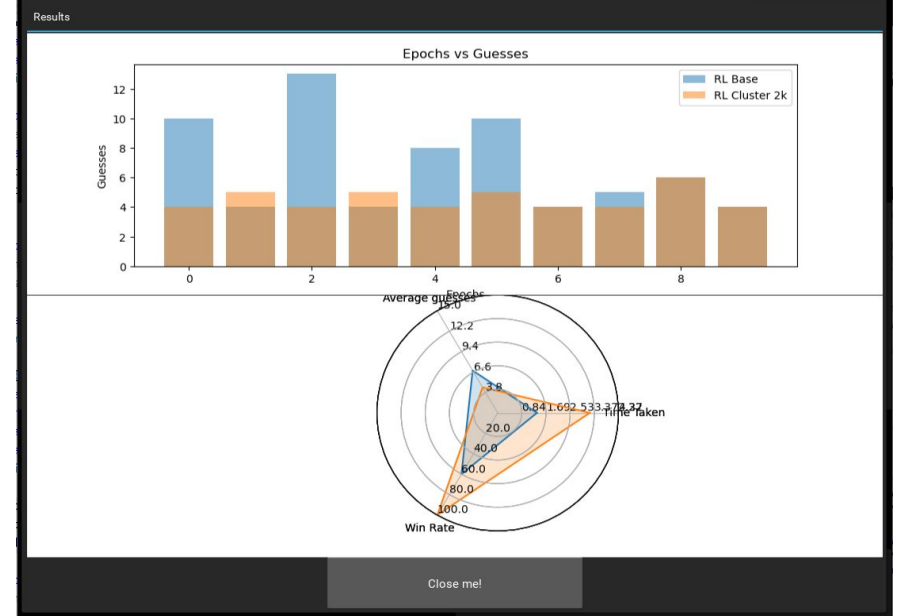
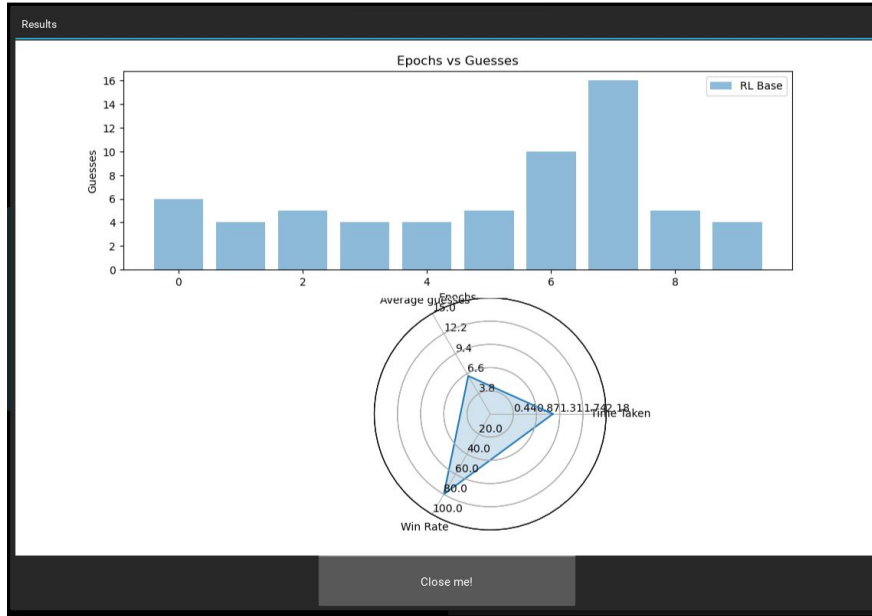
1) Kivy GUI: Customer Parameter Values

Results from grid search used as default settings in our GUI

- User can try out different settings and observe the performance



1) Kivy GUI: Visualisation & Comparison



Two screenshot of results page, with 1 run and 2 runs (for comparison) respectively

2) PyGame GUI

- Observe an agent solving the daily Wordle in hard-mode
- Keep pressing “enter” to play
- Press “esc” to rerun the solver



Screenshot of PyGame GUI Daily Wordle Solver

Wordle

A 5x4 grid of colored squares representing Wordle feedback. The colors are: Row 1: Light Gray, Yellow, Light Gray, Yellow, Yellow; Row 2: Green, Light Gray, Yellow, Light Gray, Yellow; Row 3: Green, Yellow, Light Gray, Yellow, Green; Row 4: Green, Green, Green, Green, Green.

Improvements

Possible Improvements

- Deep Q-learning
 - Might be an overkill given already how well our model performs, ~97% win-rate
- Discretization of words within clusters
 - To form a probability distribution within each cluster to select the next guess from
- Generalize for other wordle variations
 - E.g. 6 letters

Wordle



Demonstration

Wordle

Thank You!

QnA